

# Course Outline

- 1 LangGraph Orientation
- 2 LangGraph foundations
- 3 Build an Application

# LangGraph Orientation

# LangGraph

Agents and LLM applications have these challenges

- **Latency in the seconds vs ms**
  - Parallelization – to save actual latency
  - Streaming – to save perceived latency

<https://blog.langchain.com/building-langgraph>



# LangChain

LangGraph: StateGraph Essentials

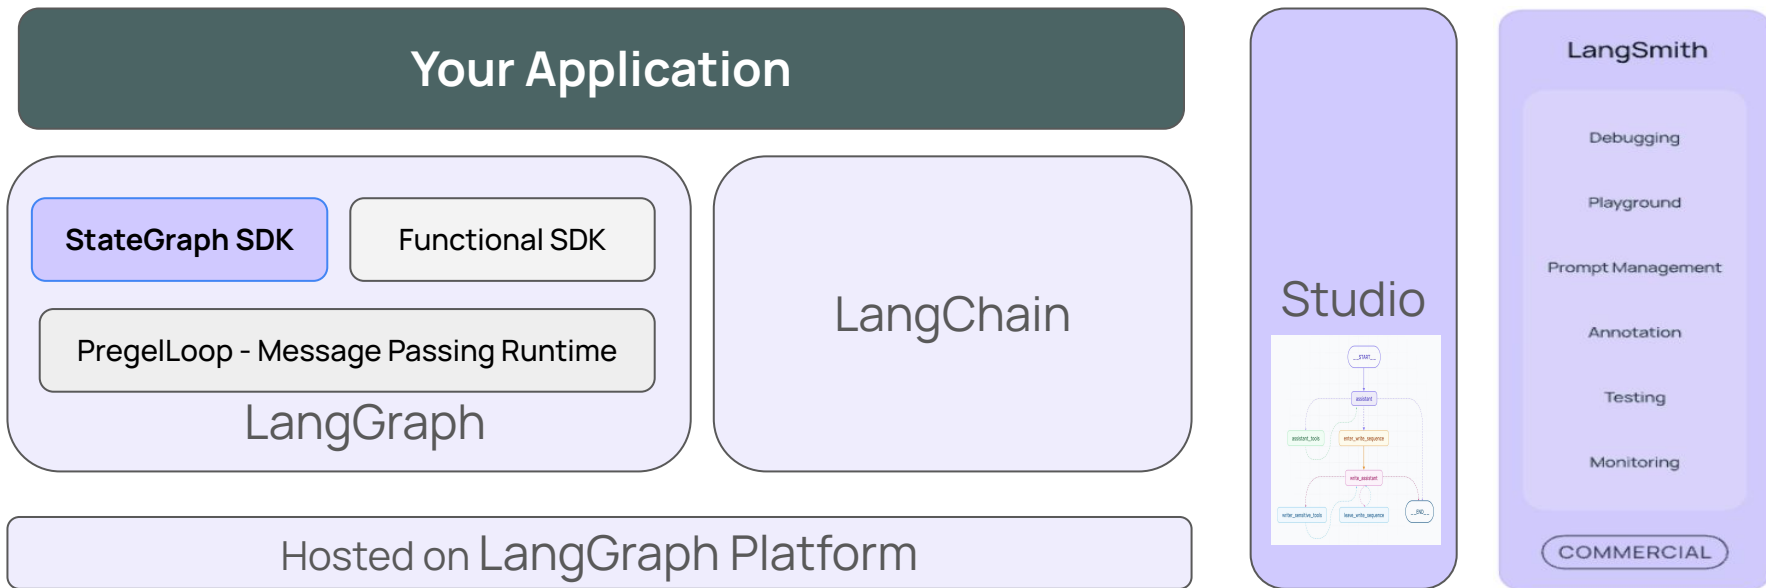


## LangGraph

[NAME]  
[TITLE], LangChain

- Parallelization – to save actual latency
- Streaming – to save perceived latency
- Task queue – to reduce number of retries (LangGraph Platform)
- Checkpointing – to reduce the cost of each retry
- Human-in-the-loop - to collaborate with the user
- Tracing - to learn how users use it (LangSmith)

# Layer Diagram



# Hidden

## Execution algorithm

Once you make the choice to structure agents into multiple discrete steps, you need to choose some algorithm to organize its execution. Even if it's a naive one that feels like "no algorithm," which is where LangGraph started before launch. The problem with using "no algorithm" is, while it may seem simpler, you're making it up as you go along, and end up with unexpected results. (For instance, an early version of a precursor to LangGraph suffered from non-deterministic behavior with concurrent nodes). The usual DAG algorithms (topological sort and friends) are out of the picture, given we need loops. We ended up building on top of the **BSP/ Pregel** algorithm, because it provides deterministic concurrency, with full support for loops (cycles).

Our execution algorithm works like this:

- Channels** contain data (any Python/JS data type), and have a name and current version (a monotonically increasing string)

- Nodes** are functions to run, which subscribe to one or more channels, and run whenever they change

- One or more channels are mapped to **input**, ie. the starting input to the agent is written to those channels, and therefore triggers any nodes subscribed to them

- One or more channels are mapped to **output**, ie. the return value of the agent is the value of those channels when execution halts

The execution proceeds in a loop, where each iteration

- Selects the 1 or more nodes to run, by comparing current channel versions and the last versions seen by each of their subscribers

- Executes those nodes in parallel, with independent copies of the channel values (ie. the state, so they don't affect each other while running)

- Nodes modify their local copy of the state while running

- Once all nodes finish, the updates from each copy of the state are applied to their respective channels, in a deterministic order (this is what guarantees no data races), and the channel versions are bumped

The execution loop stops when there are no more nodes to run (ie. after comparing channels with their subscriptions we find all nodes have seen the most recent version of their subscribed channels), or when we run out of iteration steps (a constant the developer can set).



# LangGraph: StateGraph Foundations




# LangGraph:StateGraph

## Components and Capabilities

- State: Data
- Node: Functions
- Edges: Control Flow
  - Serial, Parallel
  - Conditional
- Checkpointing/Memory
- Human In the Loop: Interrupts

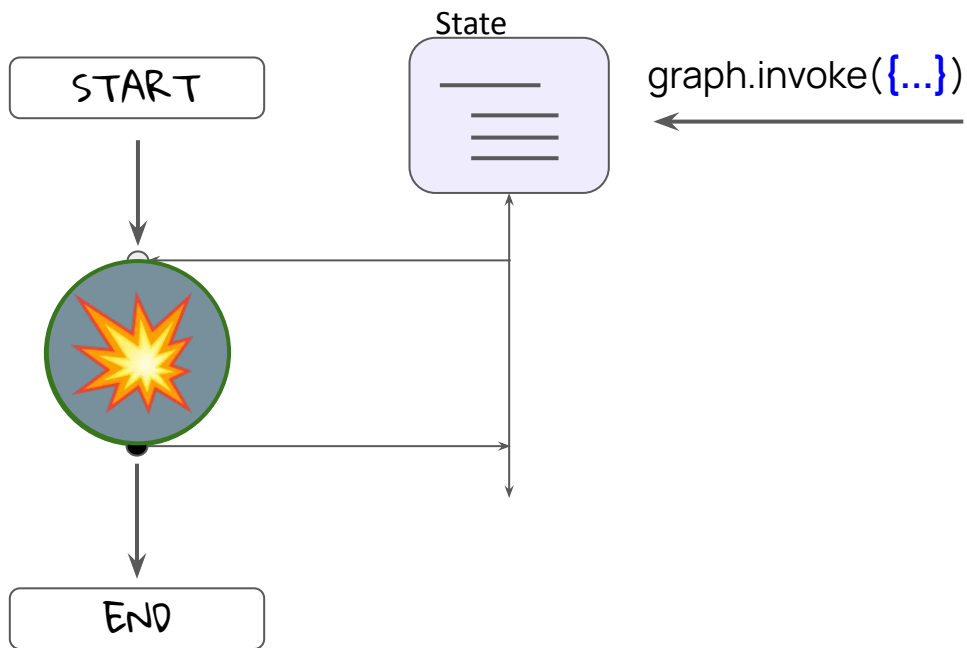
# State, Nodes



```
class State(TypedDict):  
    nlist : List[str]
```

```
def node_a(state: State):  
    ...  
    return({"nlist": [note]})
```

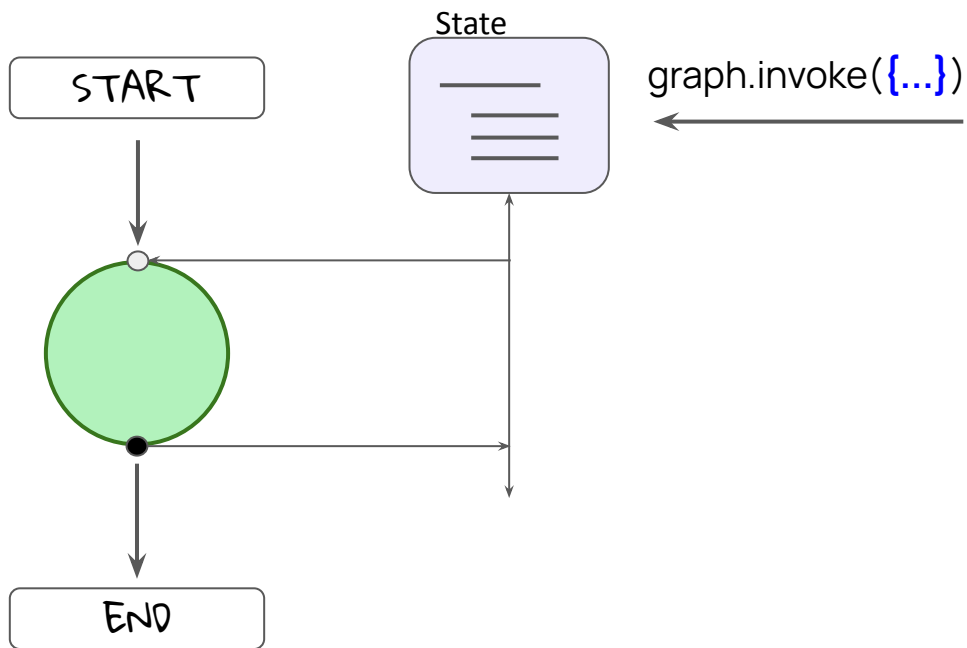
# State, Nodes



```
class State(TypedDict):  
    nlist : List[str]
```

```
def node_a (state: State):  
    ...  
    return({"nlist": [note]})
```


# State, Nodes



```
class State(TypedDict):  
    nlist : List[str]
```

```
def a (state: State):  
    ...  
    return({"nlist": [note]})
```

# State, Nodes




```
class State(TypedDict):  
    nlist : List[str]
```

```
def a (state: State):  
    ...  
    return({"nlist": [note]})
```


# Edges: Control Flow

Edge

Serial




Parallel




Conditional Edge


Conditional



Map-Reduce




# Reducers



```
class State(TypedDict):  
    nlist : List[str]
```

# Reducers



```
class State(TypedDict):  
    nlist : Annotated[list[str], operator.add]
```

type


reducer function









# Lab - Parallel/Conditional Edges



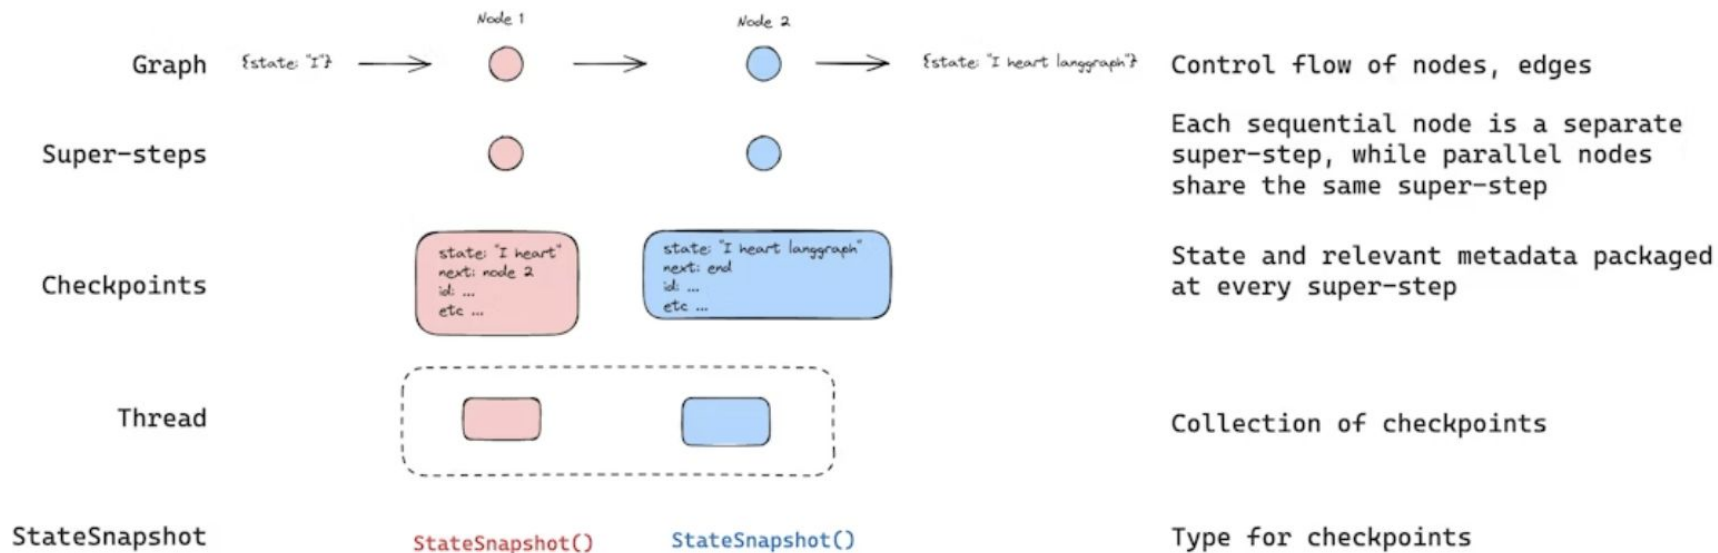
# Conditional Edge








# Memory




# Back from Lab Memory



# Memory / Checkpointers




# Memory





# Memory



# Memory

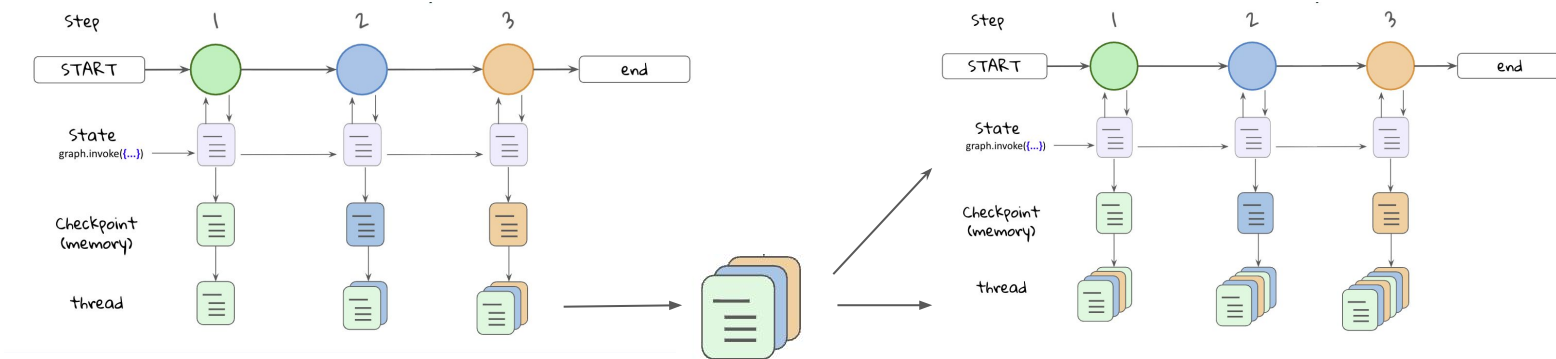
## Benefits

- **Recover gracefully from failures** — resume without losing progress.
- **Time travel** — roll back to a known good point and continue forward.

# Memory / Checkpointers

## Benefits

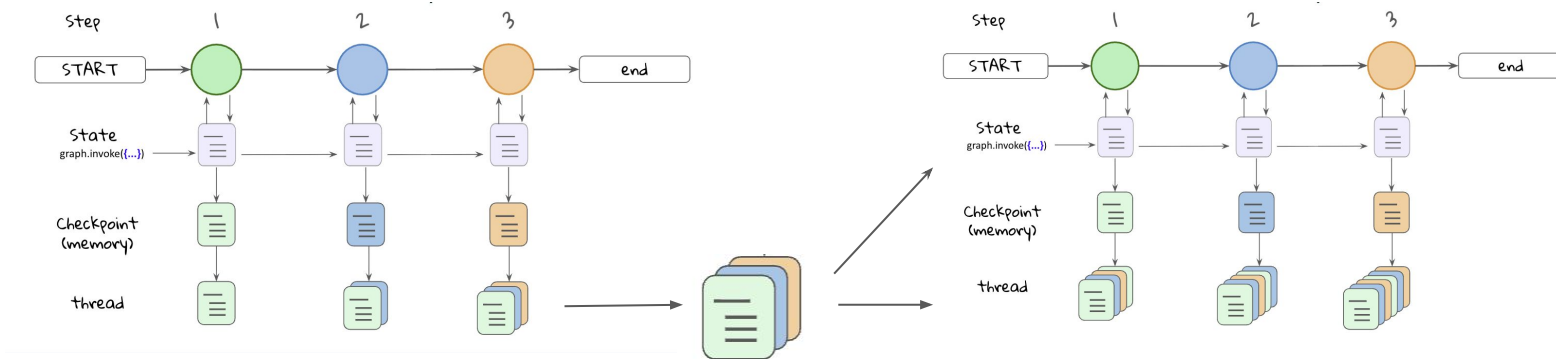
- **Recover gracefully from failures** – resume without losing progress.
- **Time travel** – roll back to a known good point and continue forward.
- **Persistent state** – data is preserved even when the graph is not running.



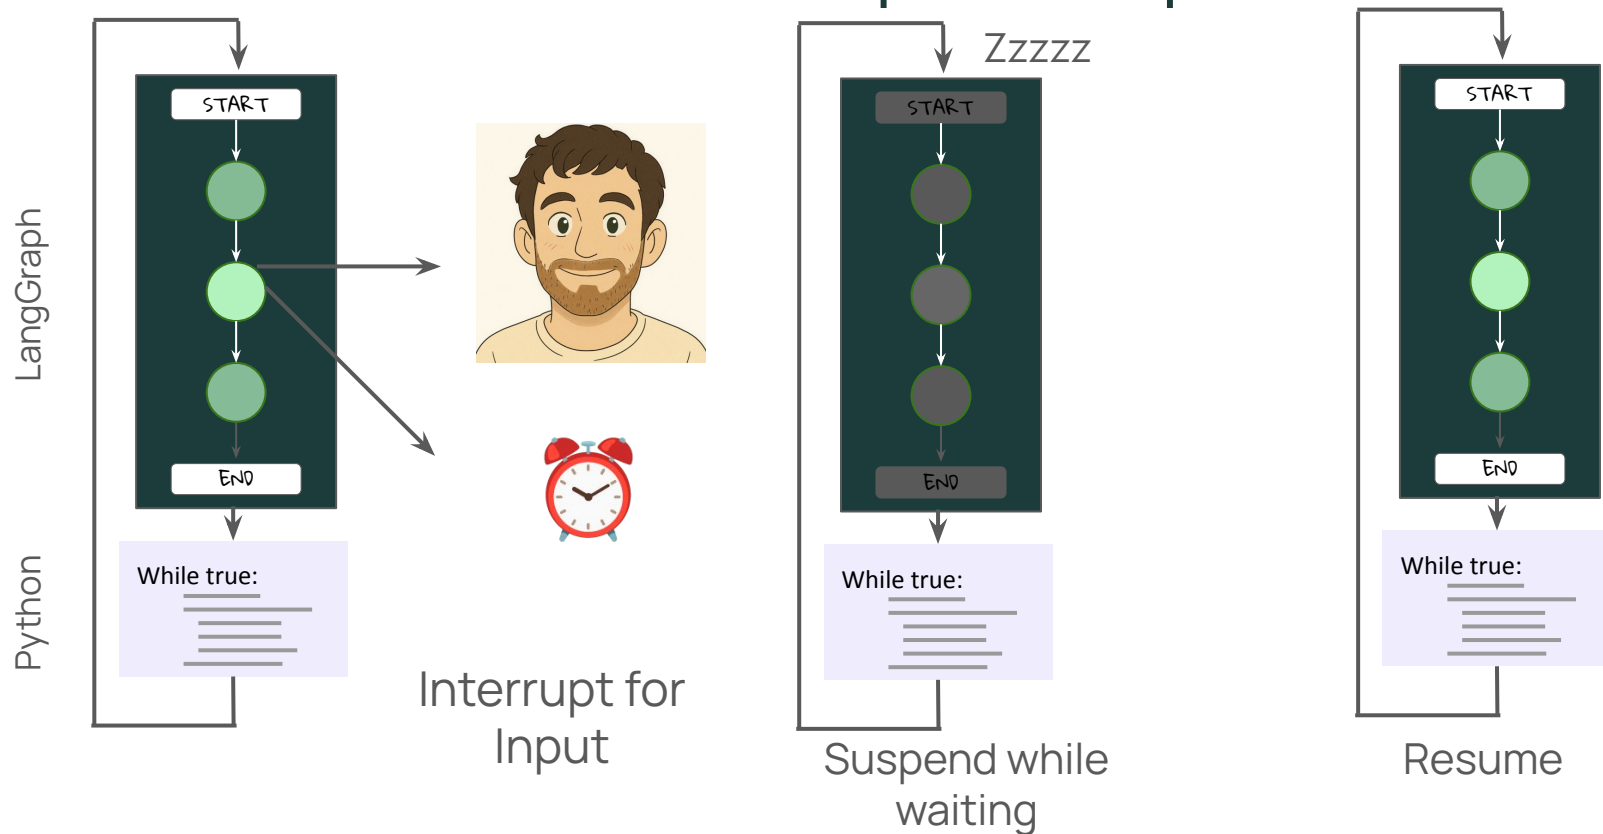
# Memory / Checkpointers

## Benefits

- **Recover gracefully from failures** – resume without losing progress.
- **Time travel** – roll back to a known good point and continue forward.
- **Persistent state** – data is preserved even when the graph is not running.
- **Restore state at any step** – pick up execution from where you left off.



# Human In the Loop: Interrupt



# Lab Interrupts



# LangGraph:StateGraph

## Components and Capabilities

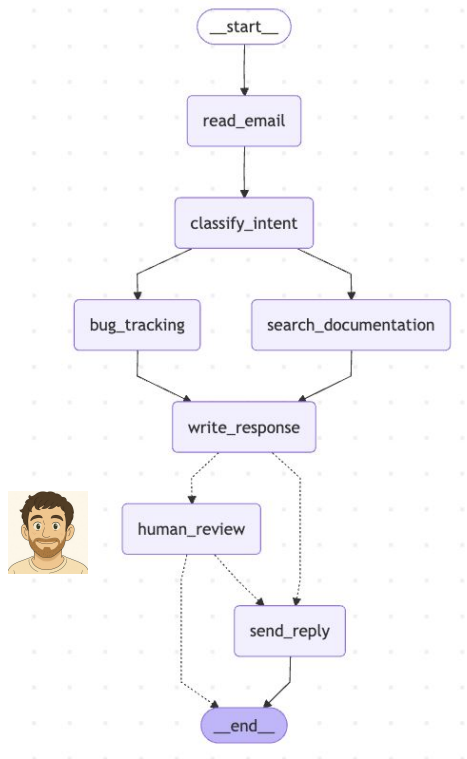
- State: Data
- Node: Functions
- Edges: Control Flow
  - Serial, Parallel
  - Conditional
- Checkpointing/Memory
- Human In the Loop: Interrupts

# Email Support Workflow


Simulate a email customer support workflow

Focus on LangGraph aspects:

- State, Nodes,
- Edges
  - Serial
  - Parallel
  - Conditional
- Memory
- Interrupt









## Edge

### Serial




### Parallel




## Conditional Edge

### Conditional



### Map-Reduce



















A solitary language model is fairly limited...




... it lacks relevant context.

Good LLM applications follow a control flow.




This control flow forms a “chain”



## Example: search → LLM (RAG) chain



Agent  $\sim$  control flow defined by an LLM




## Chain

Developer defined control flow

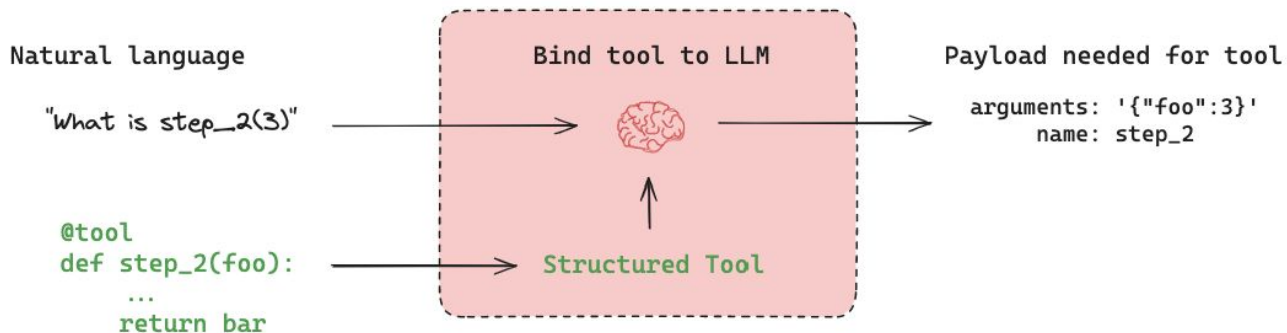


## Agent

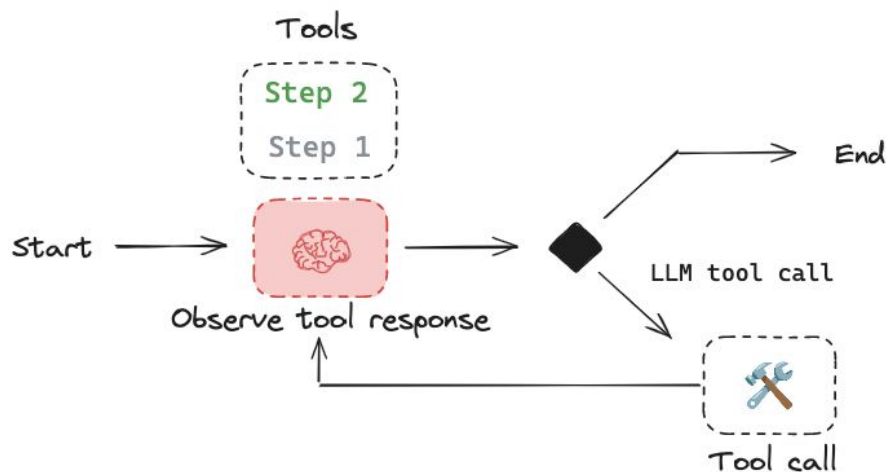
LLM defined control flow



# Agents typically use tools-calling to execute steps




## Simplest design is a `for` loop (ReAct)





# ReAct agents are flexible: any state possible!

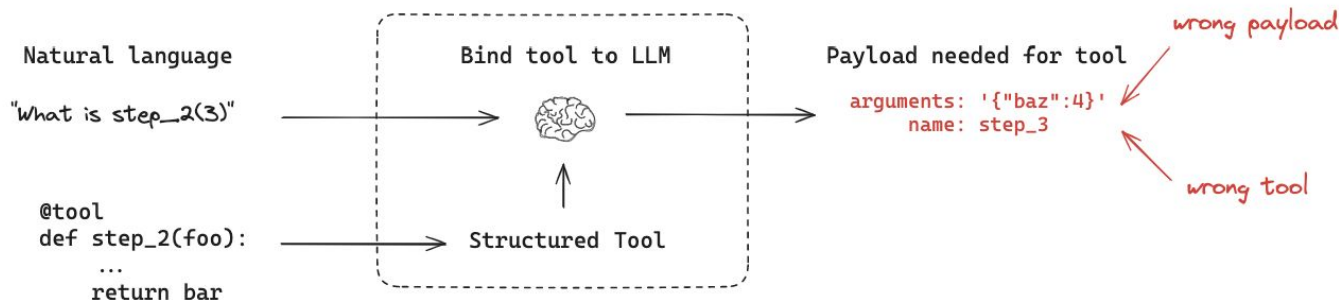



... but they suffer from poor reliability




... often caused by

- Task ambiguity
- LLM non-determinism
- Tool misuse
- Tool dependencies
- ... and more!



## Can we have both?

Chain

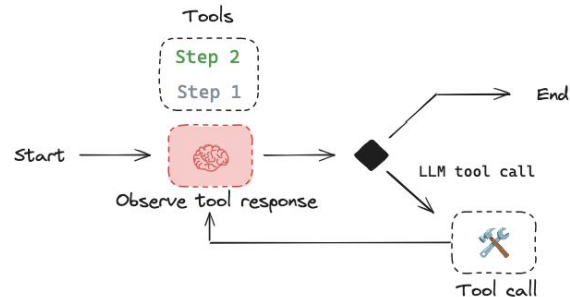


Not flexible  
More reliable

?

Flexible  
Reliable

Agent (for loop)



Flexible  
Less reliable

## 2 Introducing LangGraph

# What is LangGraph ?

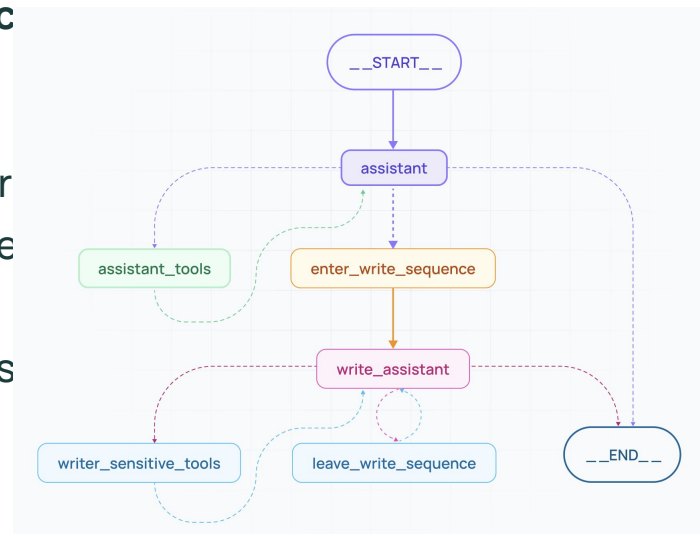
LangGraph applications balance agent control with agenc

Its core pillars support:

- A**
- B** **Controllability:** to define both explicit and implicit wor
- C** **Persistence:** to enable human-agent/multi-agent inte
- D** **Human-in-the-loop:** to facilitate human guidance
- D** **Streaming:** to expose any event (or token) as it occurs



LangGraph also:

- Works with or without LangChain
- Integrates with LangSmith




### A Controllability

Intuition: Let developer set parts of control flow (reliable)



## A Controllability


Intuition: Inject LLM to make it an agent (flexible)





### A Controllability


 **LangGraph: Express control flows as graphs**



## A Controllability

We can have both!

Chain



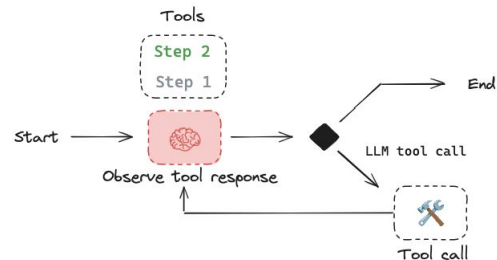
Not flexible  
More reliable

 LangGraph



Flexible  
Reliable


Agent (for loop)



Flexible  
Less reliable

## A Controllability

# LangGraph Agent



## A Controllability


LangGraph allows for developer + LLM-defined control flows

Chain



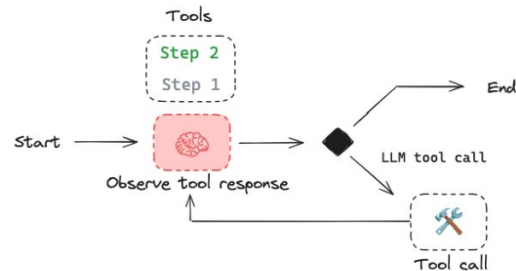
Not flexible  
Most Reliable

Agent (LangGraph)



More Flexible  
More Reliable


Agent (for loop)



Most Flexible  
Least reliable


## A Controllability

Subgraphs enable complex system design by managing states separately



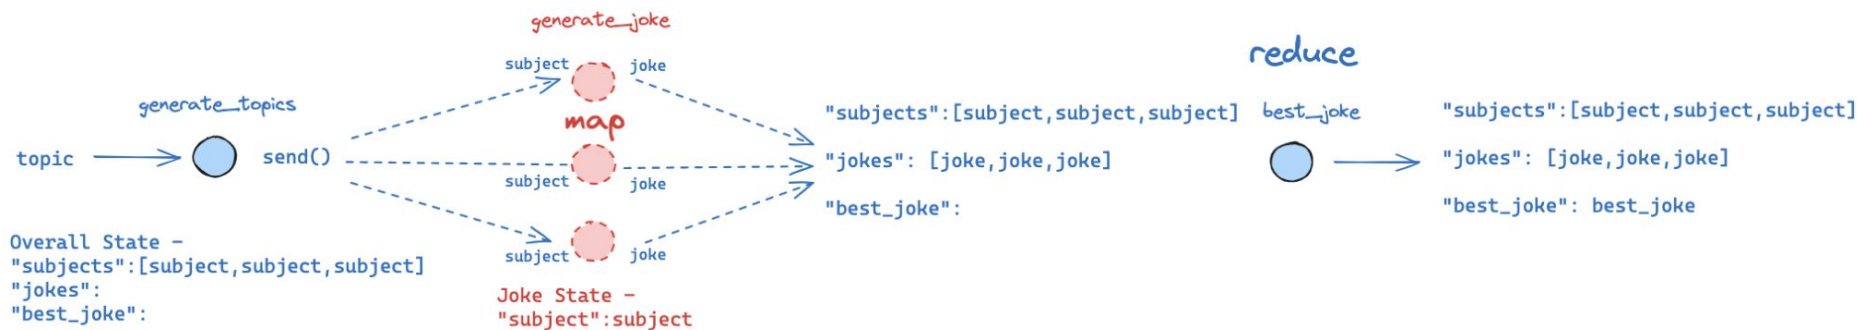
### A Controllability

Branches enable parallel execution of nodes to speed up overall graph operation




## A Controllability

Map-reduce branches enable efficient parallel processing and flexible execution



## B Persistence

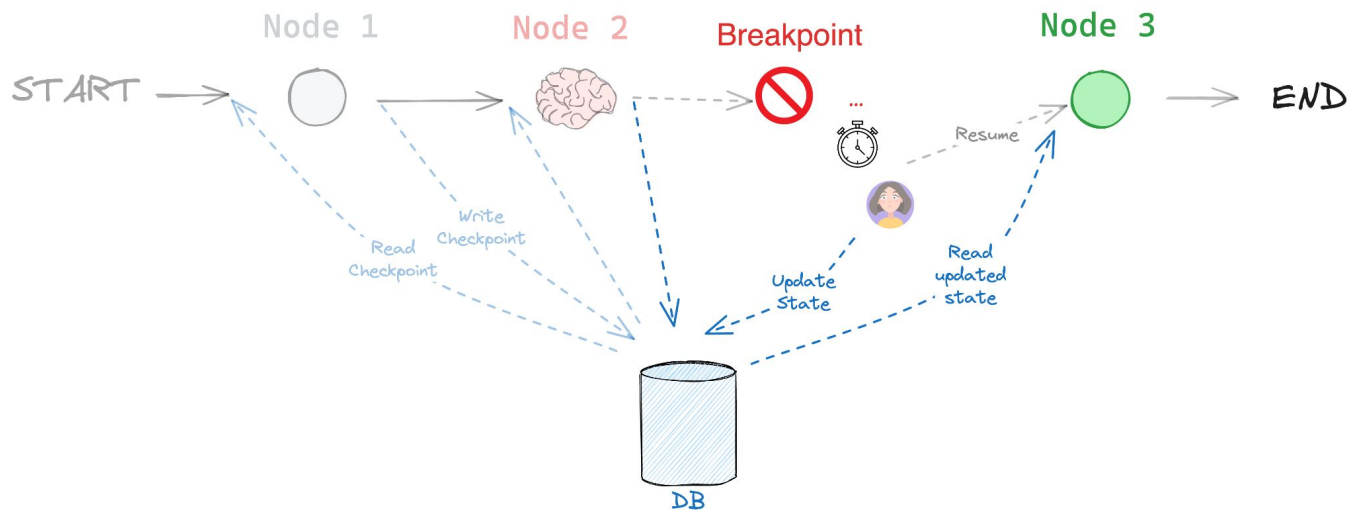
### Provides “Memory”





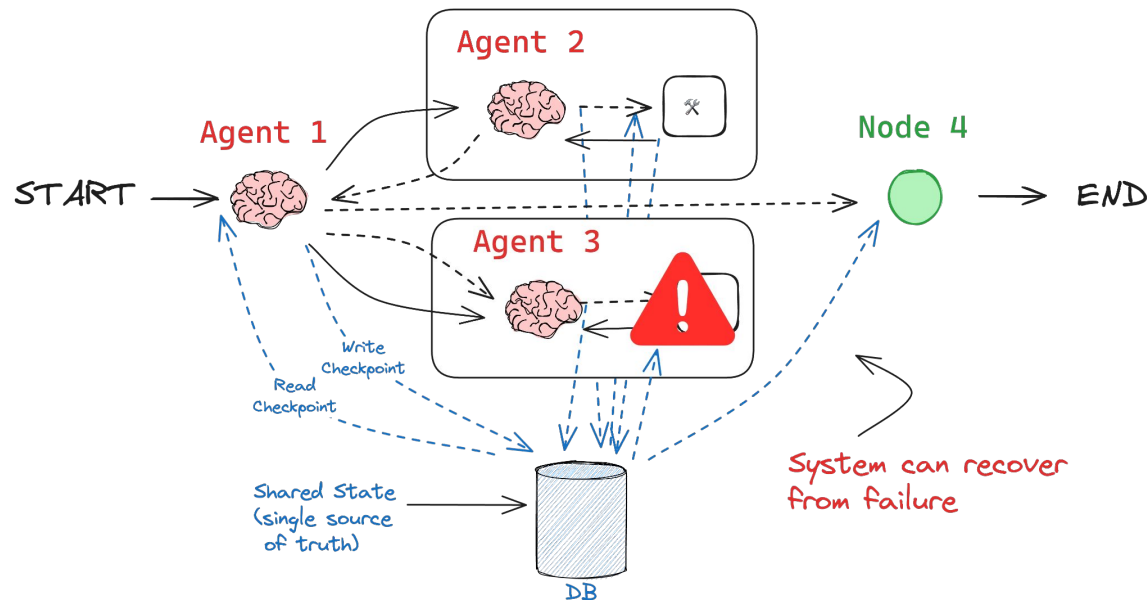
## B Persistence

# Enables human-agent interactions



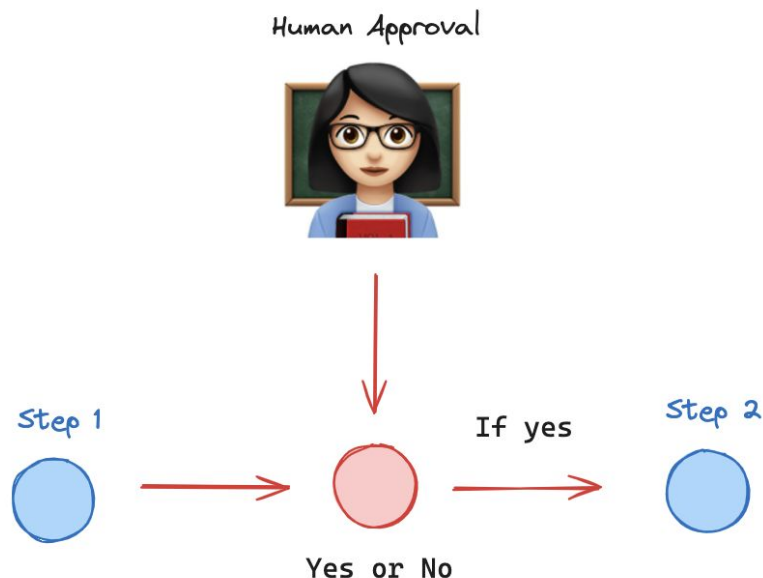
## B Persistence

Enables fault-tolerant multi-agent interactions




## c Human-in-the-loop

Breakpoints enable human-in-the-loop interactions

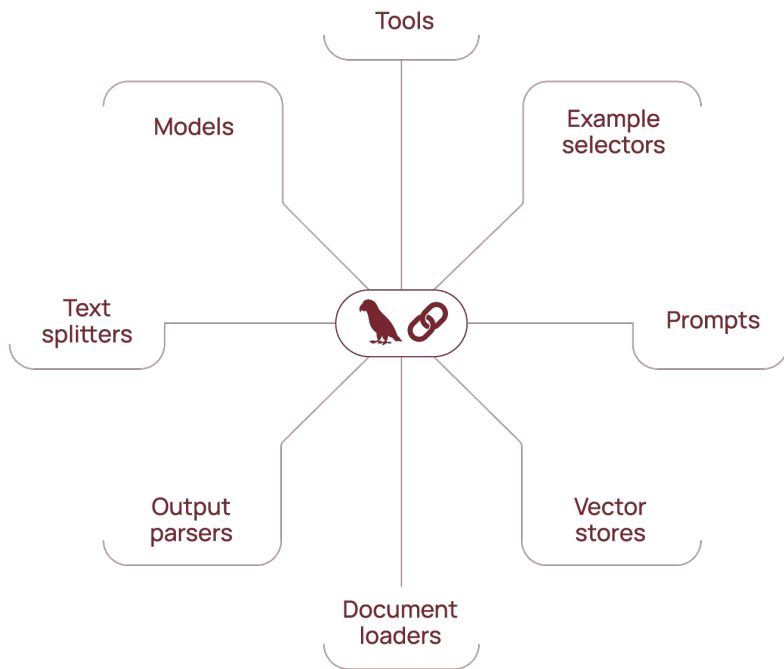


## D Streaming

First-class support for token and event-level streaming



## 3 LangGraph within LangChain's Ecosystem




LangChain's AI abstractions and integrations make it the **#1** choice for developers when building with GenAI

15M+  
Monthly  
Downloads

100K+  
Apps  
Powered

3K+  
Contributors



## LangSmith

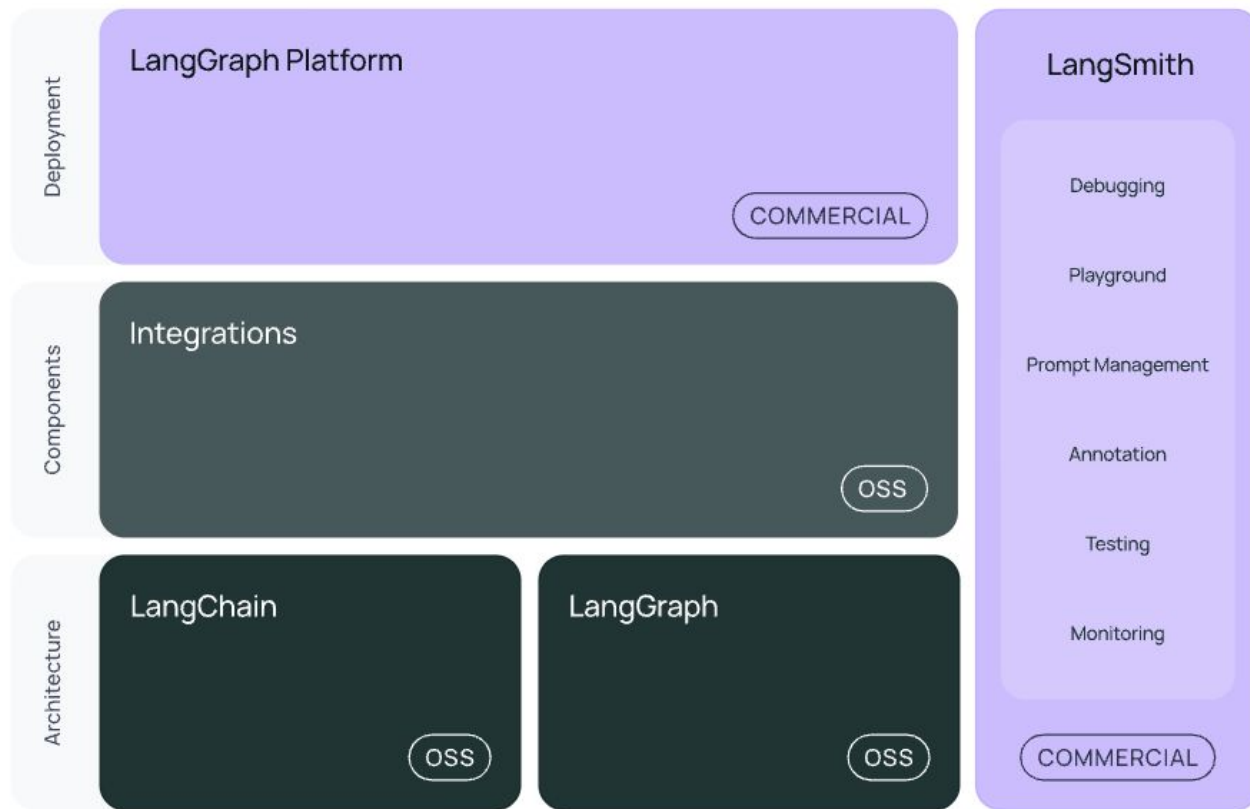
LangSmith is a unified DevOps platform, purpose-built for LLM applications

**100K+**  
Users  
signed up

**300M+**  
Traces  
logged

**30K+**  
Monthly  
active teams

### 3 LangGraph within LangChain's Ecosystem






REMOVE IF YOU DO NOT WANT  
TO INCLUDE A CODE DEMO -  
THIS IS LIKELY +15 MINUTES

## Code Demo



<https://github.com/vbarda/pandas-rag-langgraph/blob/6ebabb3c87fd4a494da7bf5c62e62aabadee4146/demo.ipynb>

# Q&A



The question: the reducer on the nodes in the superstep is applied before considering the conditional edge